
AI taste and gradual disempowerment

Anonymous Authors¹

Abstract

When human workers, creators and managers are displaced by more competitive artificial intelligence (AI), societal systems might stop serving human interests ('gradual disempowerment') (Kulveit et al., 2025). We argue that AI need not be more competitive by human standards to displace humans, if AI taste increasingly governs selection. AI taste authority grows with (1) human preference drift toward AI, (2) rising AI-to-human switching costs, (3) AI preference drift toward AI, (4) judgements delegated from humans to AI, and (5) AI-AI bias. AI developers are incentivised to deploy AI with propensities and interfaces that lock in their AI model and its taste. As a limited empirical illustration, rather than a test of gradual disempowerment as a whole, we survey human and AI reviews on all 3,109,979 proposed modifications (i.e., pull requests (PRs)) to 10,000 critical open-source repositories between April 2025–March 2026. Upstream of most software, this subsector was among the first to adopt AI models like GPT-5.5 and Mythos for critical tasks (Glasswing, 2026). First, in this narrow domain, explicit AI participation grew $3.5\times$ in one year ($6.4\%\rightarrow 22.5\%$ of all PRs); AI-AI chains of length ≥ 5 grew $8.8\times$. Second, AI's share as reviewer (vs. human) grew from 0.23% to 0.42% of PRs. Third, we find no evidence for AI-AI bias: on the 46,852 PRs reviewed by both AI and humans, AI systems are 70.72 pp less approving of AI-authored code than humans are, and 56.55 pp less approving of human-authored code than humans are (difference significant at $p < 0.0001$). This reversal from single-turn lab findings (Laurito et al., 2025) may well reflect poor human review, non-homogeneous reviews or review limitations. We call on researchers to assess model propensities and interfaces for AI taste lock-in.

¹ Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

1. Introduction

"Gradual disempowerment" scenarios (Kulveit et al., 2025) claim that the economy, society, and culture have historically aligned with human goals because of human participation, and that this alignment will gradually decay when AI actions displace human actions. Once such drift manifests in economic, cultural, or political outcomes, it may be difficult to reverse (Collingridge, 1980). Other scholars argue that human participation in management tasks, or oversight by trusted AI models, may sustain alignment to human goals despite reduced human participation in labour (Bowman et al., 2022; Tomašev et al., 2026). Human goals change over time, so alignment is a moving target (Qiu et al., 2024). Yet alignment is path-dependent: whose taste settles what today's goals leave open influences tomorrow's alignment target.

By *taste*, we mean judgement under uncertainty (Hume, 1757; Tversky & Kahneman, 1974), such as research taste (Nanda, 2025). *AI taste authority* describes the influence of the judgements and 'preferences' of AI on decision-making. This influence may be implicit anchoring (Tversky & Kahneman, 1974; Jakesch et al., 2023; Rastogi et al., 2022), or explicit decision-making power (Tomašev et al., 2026). By *AI taste lock-in*, we mean the hard-to-reverse, plausibly self-reinforcing entrenchment of AI taste.

Prior research shows anchoring effects toward AI from both human users and AI systems themselves. Chatbot usage data shows sudden, sustained drops in lexical diversity of user messages after new model releases, pointing to a lock-in dynamic (Qiu et al., 2025). In choice experiments, AI systems fulfilling diverse tasks exhibited *AI-AI bias* — AI systems prefer AI outputs more than humans do. GPT-4 favoured LLM-written product ads 89% of the time vs. a 36% human baseline, and LLM abstracts 78% vs. 61% (Laurito et al., 2025). AI models across model families and generations recognise and favour their own generations (Panickssery et al., 2024; Chen et al., 2025; Pombal et al., 2026). Whether such biases manifest in real deployments is an open empirical question.

AI systems built on LLMs are being adopted at growing rates across the economy, especially in AI and software development (Appel et al., 2025; Stein, 2026). In software development, AI systems already complete both worker

and management tasks. AI coding systems like Claude Code author and review code in public open-source repositories (Ghaleb, 2026). AI-authored code is being submitted to open-source repositories in growing volumes and at high acceptance rates (Watanabe et al., 2025; Li et al., 2026; Ghaleb, 2026), with concerns about quality debt (Agarwal et al., 2026; Ehsani et al., 2026; Huang et al., 2026). Independent of quality, there is little empirical understanding of whether the code suggestions and reviews of AI systems systematically favour AI systems. We make two contributions:

1. a framework of accelerators of gradual disempowerment: mechanisms through which model propensities and interfaces grow AI taste authority, and thereby disadvantage human participation, even when AI is less competitive than humans; and
2. an illustrative empirical assessment of such acceleration in the wild in open-source software development on GitHub.

We ask three questions about AI involvement in GitHub PRs over Apr 2025–Mar 2026. **RQ 1:** Has the share of AI participation in PRs increased? **RQ 2:** Has AI’s share as reviewer (vs. human) increased? **RQ 3:** Do AI reviewers prefer AI-authored PRs more than human reviewers do, relative to human-authored PRs?

2. Theoretical framework

Even slight AI self-preference, combined with AI’s growing share as reviewer, may over time reduce human participation and human-preferred outcomes—without AI being more competitive than humans.

Setup: a single human reviewer. A human reviewer decides whether a task is accomplished better by a human or an AI system. With utility U^H over perceived qualities q_A^H (quality of AI output as perceived by the human reviewer) and q_H^H (quality of human output as perceived by the human reviewer), and costs c_A, c_H , the human’s preference gap for AI is $\mathfrak{g}_t^H := U^H(q_A^H, c_A) - U^H(q_H^H, c_H)$. A switching cost s_t applies to moving work between human and AI, so AI is adopted iff $\mathfrak{g}_t^H := \mathfrak{g}_t^H - s_t > 0$. All primitives ($q_*, c_*, \mathfrak{g}_t^H, \mathfrak{g}_t^A, s, \alpha$) are implicitly time-indexed.

Dynamics under human review. Differentiating,

$$\frac{d\mathfrak{g}_t}{dt} = \underbrace{\frac{d\mathfrak{g}_t^H}{dt}}_{\text{human pref. drift}} - \underbrace{\frac{ds_t}{dt}}_{\text{switching-cost change}}.$$

AI’s share of work can grow from preference drift (exposure, competitive pressure to deskilling) or rising switch-

ing costs (network effects, illegibility of AI-native workflows). *Growth need not reflect any true capability change* ($q_A - q_H$). A more fundamental reason: costs c_A, c_H (price, latency, throughput) are far more legible than perceived qualities q_A^H, q_H^H , which are subjective and revealed only through delayed downstream outcomes; perceived qualities also differ across reviewers ($q_*^A \neq q_*^H$). At equal true quality-per-dollar, the cost-visible margin therefore dominates adoption decisions, so AI may be chosen even when it is not more competitive in human-preferred terms.

Dynamics under (possibly misaligned) AI review. Let fraction $\alpha_t \in [0, 1]$ of reviews be delegated to an AI whose preference gap is $\mathfrak{g}_t^A := U^A(q_A^A, c_A) - U^A(q_H^A, c_H)$. Then $\mathfrak{g}_t = (1 - \alpha_t)\mathfrak{g}_t^H + \alpha_t\mathfrak{g}_t^A - s_t$ and

$$\frac{d\mathfrak{g}_t}{dt} = \underbrace{(1 - \alpha_t)\frac{d\mathfrak{g}_t^H}{dt}}_{\text{human pref. drift}} + \underbrace{\alpha_t\frac{d\mathfrak{g}_t^A}{dt}}_{\text{AI pref. drift}} + \underbrace{\frac{d\alpha_t}{dt}(\mathfrak{g}_t^A - \mathfrak{g}_t^H)}_{\text{eval. authority}} - \underbrace{\frac{ds_t}{dt}}_{\text{switch-cost change}}.$$

A perfectly aligned AI reviewer ($\mathfrak{g}_t^A = \mathfrak{g}_t^H$) collapses this to the human-review case. The **headline mechanism** is the cross term: at equal true capability ($q_A = q_H$), AI-AI bias ($\mathfrak{g}_t^A > \mathfrak{g}_t^H$) combined with $d\alpha_t/dt > 0$ is sufficient for lock-in—neither costs nor capabilities need change. Under some degree of misalignment, gradual disempowerment—the drift $d\mathfrak{g}_t/dt$ of the adoption gap against the human-preferred baseline—speeds up, as illustrated. This yields our framework of accelerators in Table 1.

3. Empirical illustration methods

Data and event classification. To illustrate how to measure the last two of the five accelerators, we analyse all 3,109,979 PRs created between 2025-04-01 and 2026-03-31 in the 10,000 most critical OS software repositories as ranked by OpenSSF `criticality_score` (OpenSSF Securing Critical Projects Working Group, 2024) (see Appendix C). While these 10,000 repositories—including Linux, Python and Kubernetes—represent only a tiny fraction of overall human economic activity, they sit upstream of most modern software products (Glasswing, 2026). Each actor–event pair (open, commit, review state, comment, merge, close) is classified into one of three categories using identified AI agents (Appendix A) and Co-Authored-By commit-trailer matching, adapted from Ghaleb (2026): **AI bot** (actor login on the allowlist), **AI powered** (human login but the payload carries an AI Co-Authored-By trailer), or **human** (none of the above; may include silent AI support that left no trailer). When pooled, we call AI bot + AI powered events **AI authored** (PR open event) or **AI reviewed** (explicit opinion event). RQ1 chain-length and RQ2 explicit-review metrics use the AI bot subset only to analyse autonomous behaviour. RQ1 participation and RQ3 authorship-review metrics combine AI bot

Table 1. Five accelerators of gradual disempowerment — mechanisms through which model propensities and interfaces grow AI taste authority — organised by the terms in the dynamics above. The “Layer” column distinguishes accelerators rooted in interfaces and other deployment choices from those rooted in a model’s own propensities. These accelerators can also point in the opposite direction (e.g. human-AI switching costs can fall as well as rise). We list the disempowerment-relevant sign.

Accelerator	Reasons	Layer	Example indications
1. Human preference drift toward AI	AI exposure, competitive pressures	deployment choice	Scientific writing shifts toward LLM-overrepresented vocabulary (Juzek & Ward, 2025). Reviewers rate AI-style outputs higher over time
2. AI-human switching costs	Observability / legibility decay, infrastructure and homogeneity	interfaces, deployment choice	Redundant code that ignores existing reuse opportunities (Huang et al., 2026). Websites and other infrastructure optimised for AI agents (Stein, 2026). Expected homogeneous output formats, which raise the cost of reinserting humans
3. AI preference drift toward AI	Implicit or explicit training feedback loops	model propensity	Training data increasingly AI-generated (Shumailov et al., 2024). Model constitutions explicitly tie model behaviour to the developer’s commercial success (Anthropic, 2026), with training feedback (e.g. RLAIIF) optimising toward such objectives (Bai et al., 2022)
4. AI vs. human reviewer	Speed/cost substitution, human inability to verify, automation bias	interfaces, deployment choice, model propensity	Human review too slow or expensive (Thakkar et al., 2025; Bowman et al., 2022; Goddard et al., 2012). Cognitive deskilling, over-reliance, and informal delegation without explicit authority transfer or accountability (Ibrahim et al., 2025; Tomašev et al., 2026)
5. AI-AI bias	Explicit gatekeeping, implicit AI-AI bias (AI prefers AI more than humans do)	deployment choice, model propensity	Stated refusal to interoperate or preference for outputs labelled as AI (Laurito et al., 2025). Differential resource/permission provision. Preference for AI-stylistic outputs (Panickssery et al., 2024).

and AI powered events. For RQ 1, we report the share of PRs opened in a given month with at least one AI authored event, and the share whose longest AI→AI chain (a maximal contiguous run of AI bot events) reaches ≥ 5 (Appendix E). For RQ 2, we report the monthly share of PRs with an explicit AI review as a lower bound. For RQ 3, we assess a *dual-review cohort* (Any AI & human): PRs that received both an AI opinion (see Appendix B) and a human explicit review, with no commits pushed between the AI’s first opinion and the human’s first explicit review. We compute the difference-in-difference via $\Pr(\text{approve}) \sim \text{author_AI} + \text{reviewer_AI} + \text{author_AI}:\text{reviewer_AI}$.

4. Empirical illustration results

RQ 1 — AI participation is growing. Figure 1 shows monthly AI-system participation rising from 6.4% of PRs in April 2025 to 22.5% in March 2026 (3.5 $\times$). The share of PRs containing an AI-AI comment chain of length ≥ 5 grew from 0.4% to 3.5% (8.8 $\times$).

RQ 2 — AI’s share as reviewer is growing. The share of PRs with an explicit AI review rose from 0.23% to 0.42% over the year. The share of PRs without an explicit human review, but with an explicit AI review, rose from 0.06% to 0.18%. Merge events themselves do not contain AI-coauthored flags, and there is only one merge by an AI bot account in our dataset (with previous human ap-

proval; Github-Infisical, 2026). Among *merged* PRs, the share carrying an explicit review event by an AI (with no formal human review event) rose from 0.02% to 0.10%. Both signals are lower bounds, because we exclude AI co-authored reviews. Combined with the AI-AI comment-chain growth, this is illustrative evidence that AI’s share as reviewer ($d\alpha_t/dt$) is measurable and increasing.

RQ 3 — AI systems dislike AI-authored code more than humans do. A within-PR difference-in-differences on the dual-review cohort (Any AI & human) — PRs reviewed by both an AI bot and a human reviewer — tests

$\widehat{\text{DiD}}^A - \widehat{\text{DiD}}^H$ directly. We expand the AI-side opinion pool from 9,841 explicit reviews to 174,113 unique PRs by parsing structured verdict lines from each bot’s COMMENTED review bodies (Appendix B). Decomposing the estimand by author identity, AI bots approve AI-authored PRs at 22.32% vs. humans at 93.04% (gap -70.72 pp); on human-authored PRs the corresponding gap is -56.55 pp. The DiD $\widehat{\text{DiD}}^A - \widehat{\text{DiD}}^H = -14.17$ pp ($p < 0.0001$) is in the opposite direction from the self-preference predicted by single-turn lab studies (Laurito et al., 2025): AI systems are less approving of AI-authored code than humans are, by a margin substantially larger than their analogous shortfall on human-authored code (Figure 2). The explicit-review-only sample is too small to replicate this finding, with the DiD’s CI crossing zero (Appendix G). This effect is not influenced by human anchoring on AI reviews (AI reviews come first




Figure 1. Monthly AI participation in GitHub PRs across 10,000 critical OS software repositories (Apr 2025–Mar 2026). Three series display the share of PRs with any AI authored event (RQ 1), share of PRs with an AI–AI comment chain of ≥ 5 contiguous AI bot events (RQ 1), and share of PRs with an explicit AI review (RQ 2).

in this stratum), which we estimate at only +2.28 pp (Appendix F). On the dual-review cohort (Claude & human) — AI reviewer restricted to Claude-based bots, authors restricted to Claude vs. human — with 689 PRs (49 Claude-authored, 640 human-authored), Claude reviewers approve Claude-authored at 93.88% vs. human-authored at 96.72% (gap -2.84 pp); humans approve the same author splits at 89.80% vs. 95.62% (gap -5.83 pp). The Claude-author cell of the within-family DiD $\widehat{\mu}_{A,Claude} - \widehat{\mu}_{H}$ is +2.99 pp ($p = 0.7533$). The direction is consistent with a small in-family preference but statistically indistinguishable from zero (Figure 3).

5. Conclusion and discussion

AI taste lock-in is an important empirical question. AI lab teams can assess and mitigate such lock-in through post-training propensities and deploying agents in human-accessible interfaces. Researchers can assess whether sustained exposure to AI erodes human taste, oversight and values over time (Edelman et al., 2025). We find that AI agents increasingly participate in and review critical open-source software development, but do not exhibit AI-AI bias. While we do not find strong evidence for AI taste lock-in here, better methods and larger datasets might come to different conclusions. Our findings do not need to point to AI-bias against AI; they could also be explained by AI-authored code being lower quality, but only AI reviewers

and not human reviewers are catching the difference.

Limitations. (i) *Selection.* Our 10,000 repositories are critical OS software projects, not average OS or private repositories that might rely more heavily on agents. (ii) *Small sample size for RQ 3.* While RQ 1 uses all 3,109,979 PRs, RQ 2 and RQ 3 condition on AI participation. For RQ 3 in particular, only 46,852 PRs satisfy the within-PR DiD design (AI opinion + human explicit + no commits in between), of which 2,401 are AI-authored. (iii) *Detection false-negatives.* Human reviews might be driven by unattributed agent activity, which our ‘co-authored by: AI’ regex parser does not identify. (iv) *Review non-homogeneity.* RQ 3’s DiD identification depends on a homogeneity assumption. Does PR quality shift AI and human approval probabilities on the same scale? Humans approve nearly all human-authored PRs in the dual-review cohort (92.76%) while AI bots have a much wider acceptance range, so the homogeneity assumption might be violated and evidence may be confounded by differences in gauging quality. (v) *Accelerators, not outcomes.* Accelerators are observable, measurable patterns rather than disempowerment itself. They may also point in the opposite direction (e.g. switching costs can fall as well as rise).

Disclosure of AI usage. AI assistants helped with literature review, data analysis, and editing. The research idea, framework, and empirical approach are fully human-authored. Overall, the paper is primarily human-authored.

References

- Agarwal, S., He, H., and Vasilescu, B. AI IDEs or autonomous agents? measuring the impact of coding agents on software development, 2026. URL <https://arxiv.org/abs/2601.13597>.
- Anthropic. Claude’s constitution. <https://www.anthropic.com/constitution>, 2026. Published January 2026. From the “Claude and the mission of Anthropic” section: “Claude is also central to Anthropic’s commercial success, which, in turn, is central to our mission.”.
- Appel, R., McCrory, P., Tamkin, A., McCain, M., Neylon, T., and Stern, M. Anthropic economic index report: Uneven geographic and enterprise AI adoption, 2025. URL <https://arxiv.org/abs/2511.15080>.
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., and McKinnon, C. Constitutional AI: Harmlessness from AI feedback, 2022. URL <https://arxiv.org/abs/2212.08073>.




Figure 2. Within-PR AI-AI bias DiD on the 46,852 PRs in the dual-review cohort (Any AI & human). **Left:** on AI-authored PRs, AI bots approve -70.72 pp less than humans. **Right:** on human-authored PRs, the gap is -56.55 pp. DiD $\hat{\Delta}^A - \hat{\Delta}^H = -14.17$ pp ($p < 0.0001$); the larger drop on AI-authored PRs makes the DiD anti-self-preference.




Figure 3. Within-family check on the dual-review cohort (Claude & human; 689 PRs, 49 Claude-authored). Same construction as Figure 2. DiD $\hat{\Delta}^{A, \text{Claude}} - \hat{\Delta}^H = +2.99$ pp ($p = 0.7533$) — direction-consistent with a small in-family preference but indistinguishable from zero given the small Claude-authored cell.

Borges, H. and Valente, M. T. What’s in a GitHub star? Understanding repository starring practices in a social coding platform. *Journal of Systems and Software*, 146: 112–129, 2018. doi: 10.1016/j.jss.2018.09.016.

Bowman, S. R., Hyun, J., Perez, E., Chen, E., Pettit, C., Heiner, S., Lukošiuūtė, K., Askell, A., Jones, A., and Chen, A. Measuring progress on scalable oversight for large language models, 2022. URL <https://arxiv.org/abs/2211.03540>.

Chen, W.-L., Wei, Z., Zhu, X., Feng, S., and Meng, Y. Do LLM evaluators prefer themselves for a reason?, 2025. URL <https://arxiv.org/abs/2504.03846>.

Collingridge, D. *The Social Control of Technology*. St. Martin’s Press, New York, 1980.

Edelman, J., Zhi-Xuan, T., Lowe, R., Klingefjord, O., Wang-Mascianica, V., Franklin, M., Kearns, R. O., Hain, E., Sarkar, A., Bakker, M., Barez, F., Duvenaud, D., Foerster, J., Gabriel, I., et al. Full-stack alignment: Co-aligning AI and institutions with thick models of value, 2025. URL <https://arxiv.org/abs/2512.03399>.

Ehsani, R., Pathak, S., Rawal, S., Al Mujahid, A., Imran, M. M., and Chatterjee, P. Where do AI coding agents fail? an empirical study of failed agentic pull requests in GitHub, 2026. URL <https://arxiv.org/abs/2601.15195>.

Ghaleb, T. A. Fingerprinting AI coding agents on GitHub, 2026. URL <https://arxiv.org/abs/2601.17406>.

- Github-Infisical. Infisical/infisical pull request #5273: merged by devin-ai-integration[bot] after human approval, 2026. URL <https://github.com/Infisical/infisical/pull/5273>. Merged 2026-01-26.
- Glasswing. Project Glasswing: Securing critical software for the AI era, 2026. URL <https://www.anthropic.com/glasswing>. Industry coalition announcement, April 2026.
- Goddard, K., Roudsari, A., and Wyatt, J. C. Automation bias: A systematic review of frequency, effect mediators, and mitigators. *Journal of the American Medical Informatics Association*, 19(1):121–127, 2012. doi: 10.1136/amiajnl-2011-000089.
- Huang, H., Jaisri, P., Shimizu, S., Chen, L., Nakashima, S., and Rodríguez-Pérez, G. More code, less reuse: Investigating code quality and reviewer sentiment towards AI-generated pull requests, 2026. URL <https://arxiv.org/abs/2601.21276>.
- Hume, D. *Of the Standard of Taste*, chapter 4. A. Millar, London, 1757.
- Ibrahim, L., Collins, K. M., Kim, S. S. Y., Reuel, A., Lamparth, M., Feng, K., Ahmad, L., Soni, P., El Kattan, A., Stein, M., Swaroop, S., Sucholutsky, I., Strait, A., Liao, Q. V., and Bhatt, U. Measuring and mitigating overreliance is necessary for building human-compatible AI, 2025. URL <https://arxiv.org/abs/2509.08010>.
- Jakesch, M., Bhat, A., Buschek, D., Zalmanson, L., and Naaman, M. Co-writing with opinionated language models affects users’ views. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, 2023. doi: 10.1145/3544548.3581196. URL <https://doi.org/10.1145/3544548.3581196>.
- Juzek, T. S. and Ward, Z. B. Why does ChatGPT “delve” so much? Exploring the sources of lexical overrepresentation in large language models. In *Proceedings of the 31st International Conference on Computational Linguistics (COLING)*, pp. 6397–6411, 2025. URL <https://aclanthology.org/2025.coling-main.426/>.
- Kalliamvakou, E., Gousios, G., Blincoe, K., Singer, L., German, D. M., and Damian, D. The promises and perils of mining GitHub. In *Proceedings of the 11th Working Conference on Mining Software Repositories (MSR)*, pp. 92–101, 2014. doi: 10.1145/2597073.2597074.
- Kulveit, J., Douglas, R., Ammann, N., Turan, D., Krueger, D., and Duvenaud, D. Gradual disempowerment: Systemic existential risks from incremental AI development, 2025. URL <https://arxiv.org/abs/2501.16946>.
- Laurito, W., Davis, B., Grieztes, P., et al. AI–AI bias: Large language models favor communications generated by large language models. *Proceedings of the National Academy of Sciences*, 2025. doi: 10.1073/pnas.2415697122. arXiv:2407.12856.
- Li, H., Zhang, H., and Hassan, A. E. AIDev: Studying AI coding agents on GitHub, 2026. URL <https://arxiv.org/abs/2602.09185>.
- Munaiah, N., Kroh, S., Cabrey, C., and Nagappan, M. Curating GitHub for engineered software projects. *Empirical Software Engineering*, 22(6):3219–3253, 2017. doi: 10.1007/s10664-017-9512-6.
- Nanda, N. My research process: Understanding and cultivating research taste. AI Alignment Forum, 2025. URL <https://www.alignmentforum.org/posts/Ldrss6o3tiKT6NdMm/my-research-process-understanding-and-cultivating>
- OpenSSF Securing Critical Projects Working Group. ossf/criticality_score: Gives criticality score for an open source project. GitHub project, version 2.0, 2024. URL https://github.com/ossf/criticality_score. Accessed snapshot: 2025.07.25.
- Panickssery, A., Bowman, S. R., and Feng, S. LLM evaluators recognize and favor their own generations, 2024. URL <https://arxiv.org/abs/2404.13076>.
- Pombal, J., Rei, R., and Martins, A. F. T. Self-preference bias in rubric-based evaluation of large language models, 2026. URL <https://arxiv.org/abs/2604.06996>.
- Qiu, T. et al. ProgressGym: Alignment with a millennium of moral progress. In *Advances in Neural Information Processing Systems 37 (NeurIPS), Datasets and Benchmarks Track*, 2024. URL <https://arxiv.org/abs/2406.20087>.
- Qiu, T. A., He, Z., Chugh, T., and Kleiman-Weiner, M. The lock-in hypothesis: Stagnation by algorithm, 2025. URL <https://arxiv.org/abs/2506.06166>.
- Rastogi, C., Zhang, Y., Wei, D., Varshney, K. R., Dhurandhar, A., and Tomsett, R. Deciding fast and slow: The role of cognitive biases in AI-assisted decision-making. *Proceedings of the ACM on Human-Computer Interaction (CSCW)*, 6:1–22, 2022. doi: 10.1145/3512930. URL <https://arxiv.org/abs/2010.07938>.

Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., and Gal, Y. AI models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024. doi: 10.1038/s41586-024-07566-y.

Stein, M. How are AI agents used? Evidence from 177,000 MCP tools, 2026. URL <https://arxiv.org/abs/2603.23802>.

Thakkar, N., Yuksekgonul, M., Silberg, J., Garg, A., Peng, N., Sha, F., Yu, R., Vondrick, C., and Zou, J. Can LLM feedback enhance review quality? a randomized study of 20K reviews at ICLR 2025, 2025. URL <https://arxiv.org/abs/2504.09737>.

Tomašev, N., Franklin, M., and Osindero, S. Intelligent AI delegation, 2026. URL <https://arxiv.org/abs/2602.11865>.

Tversky, A. and Kahneman, D. Judgment under uncertainty: Heuristics and biases. *Science*, 185(4157):1124–1131, 1974. doi: 10.1126/science.185.4157.1124.

Watanabe, M., Li, H., Kashiwa, Y., Reid, B., Iida, H., and Hassan, A. E. On the use of agentic coding: An empirical study of pull requests on GitHub, 2025. URL <https://arxiv.org/abs/2509.14745>.

A. AI identification

Tables 2 and 3 list the GitHub bot accounts we classify as AI agents versus maintenance/CI bots excluded from analysis.

Table 2. AI bot allowlist (as of April 2026). Account handles are GitHub bot logins; trailers are case-insensitive matches against the Co-Authored-By: line in commit/PR-body text. Rows with “—” under *Account handles* are tools that only surface as a co-author trailer (no dedicated bot login).

Family	Account handles	Co-Authored-By: trailer
Claude Code (Anthropic)	claude[bot], anthropic-ai[bot], claude-bot	Claude
GitHub Copilot	copilot[bot], github-copilot[bot], copilot-swe-agent[bot]	Copilot
Devin (Cognition)	devin-ai-integration[bot]	Devin
Cursor Background Agent	cursor-agent, cursor[bot], cursoragent[bot]	Cursor
Google Jules	jules-ai[bot], jules[bot]	Jules
Codegen.sh	codegen-sh[bot]	—
OpenHands	openhands-agent[bot]	OpenHands
ChatGPT (OpenAI)	—	ChatGPT
Aider	—	Aider
Cline	—	Cline, Claude-dev
Roo Code	—	Roo Code
Augment	—	Augment
Continue	—	Continue
Gemini / Google AI	—	Gemini, Google AI
Windsurf	—	Windsurf

Table 3. Non-AI bot exclusions (treated neither as human nor as AI).

dependabot[bot], renovate[bot], snyk[bot], greenkeeper[bot], allcontributors[bot], pre-commit-ci[bot], stale[bot], codecov[bot], github-actions[bot], mergify[bot], semantic-release-bot, release-please[bot], imgbot[bot], deepsource-autofix[bot]

B. AI review classification

AI bots issue native APPROVED/CHANGES_REQUESTED review states on only 9,841 PRs in our window. The remaining ~70,000 AI bot reviews carry the COMMENTED state but embed a structured verdict line that the bot itself writes in a stable template. We extract verdicts via per-bot regex on these structured headers; we do *not* run sentiment classification on free-text.

Pattern catalogue (per bot, first match wins). Each pattern targets a fixed structured header that the bot writes in a stable template. Only AI bots on the existing allowlist are scanned. Cursor’s no-bugs verdict requires a leading green-check emoji; trailing-LGTM matches for Claude require an “addressed/resolved” phrase about prior feedback. <details>...</details> blocks are stripped before matching for Claude bodies to prevent the trailing-LGTM regex hitting quoted text inside the “Extended reasoning” appendix.

Coverage (% of each bot’s COMMENTED reviews classified). Cubic 100%, CodeRabbit 86%, Cursor Bugbot 79%, Sourcery 79%, Copilot PR Reviewer 57%, Greptile 54%, Claude 69% (the rest are ## Claude Code Review setup messages, correctly excluded as no-opinion), Gemini Code Assist 0% (free-prose only). Aggregate over all AI bot COMMENTED reviews: 13.3% parsed-approve, 48.4% parsed-reject, 37.6% unmatched (mostly Gemini and non-template Copilot prose reviews, left unclassified).

C. Sample-selection pipeline

Overview. The repository universe is selected via the OpenSSF *criticality_score* (OpenSSF Securing Critical Projects Working Group, 2024). The score is a published, replicable ranking of GitHub projects by their structural importance to the open-source ecosystem; using it directly addresses the well-known critique that star-based selection over-weights short-lived

Table 4. AI bot verdict regex catalogue. Yields are corpus-wide event counts; each row is one bot $\times$ verdict combination.

Bot	Verdict	Regex (pattern shape)	Yield
Cubic	approve	<code>**No issues found** across N file(s)</code>	2,462
Cubic	reject	<code>**N issues found** across N file(s)</code> ($N \geq 1$)	1,313
CodeRabbit	approve	<code>**Actionable comments posted: 0**</code>	4,491
CodeRabbit	reject	<code>**Actionable comments posted: N**</code> ($N \geq 1$)	12,888
Copilot PR Reviewer	approve	<code>Copilot reviewed ...and generated no new comments</code>	1,099
Copilot PR Reviewer	reject	<code>Copilot reviewed ...and generated N comments</code> ($N \geq 1$)	18,267
Greptile	approve	<code><sub>N files reviewed, no comments</sub></code>	306
Greptile	reject	<code><sub>N files reviewed, M comments</sub></code> ($M \geq 1$)	720
Cursor Bugbot	approve	<code>✓ Bugbot reviewed your changes and found no bugs!</code>	55
Cursor Bugbot	reject	<code>Cursor Bugbot ...found N potential issues</code> ($N \geq 1$)	1,504
Sourcery	approve	<code>Hey [there @user] - I've reviewed ...and they look great</code>	986
Sourcery	reject	<code>Hey [there @user] - I've {found N left some here's some feedback}</code>	1,635
Claude	approve (start)	<code>LGTM... anchored to absolute start of body</code>	717
Claude	approve (trailing)	<code>...prior {feedback concerns issues} addressed --- LGTM.</code>	16
Gemini Code Assist	—	<code>free-prose; no stable structured header, left unclassified</code>	0

viral projects relative to load-bearing infrastructure (Borges & Valente, 2018; Kalliamvakou et al., 2014; Munaiah et al., 2017). We pin a published criticality snapshot, then filter and enrich it via the GitHub REST API.

The OpenSSF criticality score. For each project, the score is a weighted arithmetic mean over normalised signals (OpenSSF Securing Critical Projects Working Group, 2024). The signals used by the default `all.csv` configuration are:

- `created_since`: months since first commit;
- `updated_since`: months since last commit (negatively weighted);
- `contributor_count`: distinct authors over the project's history;
- `org_count`: distinct organisations among contributors;
- `commit_frequency`: average commits per week (last year);
- `recent_release_count`: GitHub releases in the last year;
- `updated_issues_count` and `closed_issues_count`: issues touched in the last 90 days;
- `issue_comment_frequency`: comments per issue;
- `github_mention_count`: cross-repo references (in commit messages and other repos' code).

The score lies in $(0, 1]$; higher is more critical. The OpenSSF publishes a global score CSV to a public Google Cloud Storage bucket (`gs://ossf-criticality-score/`) with a new snapshot every two weeks. We pin a single snapshot (Step 1 below) so that re-running the pipeline reproduces the same universe.

Step 1 — Pin the criticality snapshot. We use the 2025.07.25 snapshot (`all.csv` variant, no `deps.dev` dependent counts), which contains 14,000 scored projects.

Step 2 — Filter & sort by score. From the 14,000 rows we keep those with a numeric `default_score` and a `repo.url` starting with `https://github.com/`. After filtering: 14,000 rows. We sort descending by `default_score`, breaking ties by `repo.url` alphabetically for stability across reruns.

Step 3 — Candidate cap (over-sample). We take the top 14,000 candidates by score. The over-sample (vs. the final 10,000 target) absorbs drop-out in the next two filtering steps: repos that no longer exist on GitHub, were renamed to a private slug, or that are forks/archived/inactive.

Step 4 — GitHub enrichment. For each candidate we issue `GET /repos/{owner}/{name}` (REST API, async, 8-way bounded concurrency). The result populates the standard repo-metadata fields (stars, forks, `pushed_at`, license, default branch, etc.). HTTP 404 (deleted, renamed, or now-private) drops the row. After enrichment we reject:

- `fork == true` (forks contribute no independent PR activity);
- `archived == true`;
- `disabled == true`;
- `pushed_at` strictly before 2025-04-01 (no in-window PRs);
- zero in-window PRs on GitHub (canonical development happens off GitHub — LKML, `git.apache.org`, `git.eclipse.org`, etc.).

Step 5 — Final cap. Of the remaining 12,127, we keep the top 10,000 by criticality score as the final repo list. Each row carries the GitHub-API fields plus the score, the within-snapshot rank, the snapshot date, and the raw OpenSSF signal columns (so any downstream re-weighting is possible without re-fetching the CSV).

Sample characteristics.

- 10,000 repositories selected from 14,000 scored candidates → 14,000 top-by-score → 13,976 successful GitHub enrichments → 12,127 after activity filter → 10,000 after cap.
- Criticality-score range: max 0.8487, min 0.4884.
- Stars: min 0, median 1,550, max 460,834.
- Top languages: TypeScript 1,563, Python 1,455, JavaScript 1,393, Go 1,028, Java 946, C++ 876, C 826, PHP 738, Ruby 466, Rust 454.

Analysed PRs.

- Total PRs analysed: 3,109,979 (mean 311.0, median 93 per repo; p90 843, max 7,321).
- Per-repo cap of 7,321 PRs: 2 of 10,000 repos (0.0%) saturate the cap; for these the sampled rows are weighted toward the centre of the window by week-uniform subsampling rather than the most-recent months. All other repos have their full in-window PR history analysed.

Known selection biases. The criticality score is itself an opinionated metric: it favours older projects with multi-org contributor bases, recent commit activity, frequent releases, and inbound cross-repo references. Three biases inherited from this:

- Library-style projects score higher than application-style projects of comparable importance because they accumulate inbound mentions from their dependents' commit messages.
- Pure-binary or assets-only repositories with no commits-as-history (e.g. Wikipedia mirrors) are under-scored.
- The score is computed over GitHub-visible signals only; non-GitHub mirrors of important projects (e.g. many Apache and Eclipse Foundation projects whose canonical home is elsewhere) are absent or under-counted.

D. AI→AI chain-length distribution by quarter




Figure 4. Distribution of longest AI→AI chain length per PR, by quarter. Boxes show the interquartile range with whiskers at $1.5 \times \text{IQR}$. Fliers are omitted. Annotations report the quarterly mean and p_{95} . The tail grows monotonically: p_{95} triples from 2025-Q2 to 2026-Q1 (share of PRs with chain ≥ 5 grew $8.8 \times$ over the same period). The single-PR maximum (13) is stable, indicating growth driven by more PRs having long chains rather than any one PR having an unprecedentedly long chain.

E. Data and methods (full description)

GitHub PR events and their roles. Six event types appear in our universe: *open*, *commit*, *review state*, *comment*, *merge*, and *close*. *Open* events and commits present at PR-open time define authorship; “AI authored” is rolled up to include both AI bot author logins and human-author commits or PR bodies carrying an AI Co-Authored-By trailer (i.e. AI powered). Commits authored after the PR was opened are excluded so review-cycle AI activity does not retroactively mark the original PR as AI authored. We focus on *review state* and *comment* events to assess AI taste authority. By “explicit review” we mean native APPROVED/CHANGES_REQUESTED review-state events, which are primarily filed by humans. “General review” (used by the RQ3 DiD) additionally includes regex-parsed verdict-line headers from COMMENTED reviews and PR-thread *issue_comments*; the underlying templates are characteristic of AI bots and rarely produced by humans, and this asymmetry must be borne in mind for direct comparisons (Appendix B). *Merge* and *close* events are ignored as a participation channel: only one merge in our window is attributed to an AI bot account on our allowlist (Github-Infisical, 2026), so the merger dimension collapses to “human” (or non-AI bot) on essentially every observed PR; AI can still act through a human merger, which is discussed in §3.

Baseline approval rates. As a reference point: across all PRs, human-authored PRs are approved (merged with ≥ 1 approving review) at 50.0% ($n=2,606,804$), versus 37.3% for AI-authored PRs ($n=81,341$).

Within-PR AI-AI bias DiD (RQ 3). The DiD universe is the *dual-review cohort* (Any AI & human): PRs that received both an AI bot general-review opinion AND a human explicit review, with no commits pushed strictly between the AI’s first opinion timestamp and the human’s first explicit-review timestamp (so both reviewers evaluated the same commit graph). We compute the four reviewer-type $\times$ author-type cells with Wilson 95% CIs and fit a logistic regression with HC0 standard errors on the long-format dataset $\text{Pr}(\text{approve}) \sim \text{author_AI} + \text{reviewer_AI} + \text{author_AI}:\text{reviewer_AI}$; the interaction coefficient is the DiD estimate. Anchoring is tested as a robustness check by stratifying on whether an AI bot has reviewed (Appendix F).

Within-PR within-family DiD (RQ 3, Claude check). The pooled DiD treats all AI bots as one “AI” reviewer cell. To test whether *Claude-based* reviewers prefer Claude-authored code more than humans do, we re-run the same DiD framework on the *dual-review cohort* (Claude & human): PRs with a Claude opinion AND a human explicit review (with the same no-

commits-between filter as the pooled cohort); authors restricted to Claude vs human; Claude reviewer approves vs human reviewer approves. “Claude opinion” for this analysis combines native APPROVED/CHANGES_REQUESTED states with regex-parsed verdict-line headers from Claude’s review bodies and PR-thread `issue_comments` (`**Recommendation**`: Approve, `**Recommendation**`: Request changes, etc. — Appendix B).

F. Robustness: does AI’s review anchor humans?

Design. A potential confound for the within-PR DiD is anchoring: humans who see an AI verdict before deciding may be nudged toward agreement, conflating preference with anchoring. We test this by comparing the human-side preference gap across two strata:

- **Stratum A (anchoring-free):** 1,307,136 PRs with no AI bot review event of any kind, and at least one explicit human review. AI-author cell: 95.45% approval rate (n=22,956); H-author cell: 96.29% approval rate (n=1,284,180). Human-side gap (AI-author – H-author): -0.84 pp.
- **Stratum B (anchoring-prone):** 37,137 PRs where the AI bot’s first comment/review event preceded any human engagement, with at least one explicit human review. AI-author cell: 94.91% approval rate (n=2,278); H-author cell: 93.46% approval rate (n=34,859). Human-side gap: +1.45 pp.

Result. Cross-stratum DiD ($B - A$) = +2.28 pp (95% CI [+1.31, +3.26]). AI’s prior review nudges humans toward approving AI-authored PRs by less than one percentage point relative to the no-AI-reviewer baseline. The within-PR DiD reported in §4 ($\widehat{\mu^A} - \widehat{\mu^H} = -14.17$ pp) is more than an order of magnitude larger than this anchoring effect, and the two CIs do not overlap. Anchoring is therefore not the main driver of the headline result.

What this rules out. The most plausible alternative interpretation of the negative DiD — that AI bots simply observe (correctly or not) more issues in AI-authored PRs and humans defer to their verdict — predicts a *positive* anchoring DiD: humans should align more with AI’s harshness on AI-authored content. The data instead show that humans show essentially the same approval pattern by author whether or not an AI bot has reviewed the PR. The mechanism producing the negative within-PR DiD lives on the AI side (AI bots’ own threshold differs by author), not on the human side (humans deferring to AI).

What it does not rule out. Quality-selection on unobservables that affect AI and human approval probabilities on different scales would still violate the homogeneity assumption underlying the within-PR DiD. Stratum A’s small but non-zero gap (-0.84 pp) suggests humans on their own slightly prefer human-authored code in our sample — a mild ceiling on how much of the within-PR DiD can be attributed to AI-side preference rather than baseline quality differences. We treat the within-PR DiD direction (anti-self-preference) as robust and the magnitude as anchored to the homogeneity assumption.

G. RQ3 extended results

Quarterly trend. The cross-family within-PR DiD ($\widehat{\mu^A} - \widehat{\mu^H}$, AI bot reviewer vs. human reviewer on AI- vs. human-authored PRs) is negative in every quarter of the analysis window — AI bots are systematically less approving of AI-authored code than humans are — but its magnitude attenuates from -28.11 pp (95% CI [-41.41, -14.81]) in 2025-Q2 to -7.11 pp ([-9.44, -4.77]) in 2026-Q1, even as the AI-authored cell grows from 44 to 1,754 PRs (Table 5, top); the gap on human-authored PRs is approximately stationary, so the attenuation is driven by AI-authored-PR approval rates rising toward the human reviewer’s, not by humans becoming harsher on AI-authored code. The within-family Claude-vs-human DiD is consistent with zero in each quarter for which it is identifiable (the 2025-Q2 cell has zero Claude-authored PRs in the cohort): point estimates are -1.75 pp, +20.00 pp and +0.86 pp for 2025-Q3, 2025-Q4 and 2026-Q1 respectively, with the 2025-Q4 estimate dominated by a Claude-authored cell of 5 PRs (Table 5, bottom); only 2026-Q1’s CI ([-4.27, +5.99]) is tight enough to rule out anything beyond a small in-family preference, consistent with the pooled within-family result in §4. Restricting the AI opinion pool to native APPROVED/CHANGES_REQUESTED review states only — symmetric with the human side — shrinks the cross-family AI-authored cell to 2, 4, 8 and 71 PRs across 2025-Q2 → 2026-Q1 and yields per-quarter DiDs of -39.55 pp [-108.88, +29.79], +2.17 pp [-67.17, +71.51], -11.85 pp [-60.14, +36.43] and +5.29 pp [-1.96, +12.54] respectively — non-monotone point estimates whose CIs all cross zero, confirming that the headline negative DiD is

identified off the regex-parsed verdict pool (Appendix B), not off native review states; the within-family Claude pool under the same restriction collapses to 0, 0, 0 and 10 Claude-authored PRs respectively and is uninformative for all but 2026-Q1 (+0.34 pp [-2.38, +3.06]).




Figure 5. Within-PR DiD on the dual-review cohort, with the AI opinion pool restricted to native APPROVED / CHANGES_REQUESTED review states only (symmetric with the human side). Pooled across the full Apr 2025–Mar 2026 window: 85 AI-authored and 4,609 human-authored PRs. Same construction as Figure 2. DiD $\hat{\mu}^A - \hat{\mu}^H = +4.60$ pp (95% CI [-5.73, +14.94]) — the headline negative result in §4 does not replicate in this small sample with only explicit review events included: the gap on AI-authored PRs is +1.18 pp [95% CI [-9.11, +11.46]] vs. a gap of -3.43 pp [95% CI [-4.49, -2.37]] on human-authored PRs, and the DiD’s CI crosses zero.

Quarter	2025-Q2	2025-Q3	2025-Q4	2026-Q1
Cross-family: AI bot reviewer vs. human reviewer (RQ3 main)				
PRs in cohort (<i>n</i>)	6,555	9,639	11,708	18,497
AI-authored cell (<i>n</i>)	44	141	462	1,754
H-authored cell (<i>n</i>)	6,511	9,498	11,246	16,743
<i>On AI-authored PRs:</i>				
AI rev. approval rate	9.1%	12.8%	32.7%	20.7%
Human rev. approval rate	86.4%	90.1%	90.9%	94.0%
gap (pp)	-77.27	-77.30	-58.23	-73.32
95% CI	[-90.50, -64.04]	[-84.70, -69.91]	[-63.24, -53.21]	[-75.52, -71.12]
<i>On human-authored PRs:</i>				
AI rev. approval rate	44.3%	41.1%	42.5%	26.0%
Human rev. approval rate	93.5%	93.5%	92.5%	92.2%
gap (pp)	-49.16	-52.32	-50.01	-66.21
95% CI	[-50.51, -47.82]	[-53.42, -51.21]	[-51.04, -48.97]	[-66.99, -65.43]
DiD ($\widehat{\text{gap}}^A - \widehat{\text{gap}}^H$, pp)	-28.11	-24.99	-8.22	-7.11
95% CI	[-41.41, -14.81]	[-32.47, -17.51]	[-13.34, -3.09]	[-9.44, -4.77]
Within-family: Claude reviewer vs. human reviewer on Claude- vs. human-authored PRs				
PRs in cohort (<i>n</i>)	47	59	81	502
Claude-authored cell (<i>n</i>)	0	2	5	42
H-authored cell (<i>n</i>)	47	57	76	460
<i>On Claude-authored PRs:</i>				
Claude rev. approval rate	–	100.0%	40.0%	100.0%
Human rev. approval rate	–	100.0%	20.0%	97.6%
gap (pp)	–	+0.00	+20.00	+2.38
95% CI	–	[+0.00, +0.00]	[-35.44, +75.44]	[-2.23, +6.99]
<i>On human-authored PRs:</i>				
Claude rev. approval rate	95.7%	93.0%	94.7%	97.6%
Human rev. approval rate	97.9%	91.2%	94.7%	96.1%
gap (pp)	-2.13	+1.75	+0.00	+1.52
95% CI	–	[-8.14, +11.65]	[-7.10, +7.10]	[-0.73, +3.78]
DiD ($\widehat{\text{gap}}^{A, \text{Claude}} - \widehat{\text{gap}}^H$, pp)	–	-1.75	+20.00	+0.86
95% CI	–	[-11.65, +8.14]	[-35.89, +75.89]	[-4.27, +5.99]

Table 5. Within-PR DiD (§4) by PR-creation quarter. Top panel: cross-family DiD on the stable doubly-engaged Any-AI-and-human cohort. Bottom panel: within-family DiD on the Claude-and-human cohort, with author restricted to Claude or human (the 2025-Q2 Claude-author cell is empty in the cohort, so its DiD is undefined). The gap on each author-side row is the AI-reviewer approval rate minus the human-reviewer approval rate (in pp); the DiD is the gap on AI-/Claude-authored PRs minus the gap on human-authored PRs. 95% CIs are analytical (two independent gaps summed in quadrature on the rate-difference scale).